

Kickstarter Analysis Proposal

10/9/2020

The Laureates: Farzeen Najam, Olivia Olsher, Vincent Liu, Emile Therrien

Introduction

Have you ever had an ambitious idea but did not have resources to pursue it? Kickstarter is a crowdfunding platform for creators who need support for their projects. Since launching in 2009, 19 million people have pledged to back up projects, and nearly 190,000 projects have been successfully funded.

However, many more projects failed to reach their goals, and our group is interested in analyzing what made projects succeed and fail. The dataset that we use comes from Kaggle, which contains Kickstarter projects up until January 2018.

We will be exploring which variables influence the success of a Kickstarter project by observing which types of projects are more likely to be funded. Our questions include:

- Does the amount of money a project asks for influence its chance of success?
- Does the category of project influence its chance of success?

Our hypotheses include:

- The more money the project asks for, the less successful it will be in terms of getting funding.
- The category “Technology” will be the most likely funded type of projects.

Data description

This data set contains 15 variables and 378,661 observations, where each observation is one kickstarter project.

The categorical variables include the name of each project, the category of each project (music, narrative film, restaurant, etc.), a broader category of each (food, film, publishing, etc.), the crowdsourcing currency, the state of each project (failed, successful, or cancelled), and the country of origin for each project.

The numerical data are the project’s ID number, the monetary goal for each, the money pledged to each project, how many backers of each project; there are three numerical variables that are not self-explanatory, `usd_pledged`: conversion in US dollars of the pledged column (conversion done by kickstarter), `usd_pledge_real`: conversion in US dollars of the pledged column (conversion from Fixer.io API), and `usd_goal_real`: conversion in US dollars of the goal column (conversion from Fixer.io API).

There are also two date columns, one for the project launch and the other for the crowdsourcing deadline.

The data were collected from Kickstarter Platform likely using web scraping methods on their own site, to be used by data scientists to model whether or not a project will be successful or not when it is launched.

Glimpse of data

```
library(tidyverse)

kickstarter <- read_csv("data/ks-projects-201801.csv")
glimpse(kickstarter)
```

```
## Rows: 378,661
## Columns: 15
## $ ID          <dbl> 1000002330, 1000003930, 1000004038, 1000007540, 10...
## $ name        <chr> "The Songs of Adelaide & Abullah", "Greeting From ...
## $ category     <chr> "Poetry", "Narrative Film", "Narrative Film", "Mus...
## $ main_category <chr> "Publishing", "Film & Video", "Film & Video", "Mus...
## $ currency     <chr> "GBP", "USD", "USD", "USD", "USD", "USD", "USD", "...
## $ deadline     <date> 2015-10-09, 2017-11-01, 2013-02-26, 2012-04-16, 2...
## $ goal         <dbl> 1000, 30000, 45000, 5000, 19500, 50000, 1000, 2500...
## $ launched     <dtm> 2015-08-11 12:12:28, 2017-09-02 04:43:57, 2013-01...
## $ pledged      <dbl> 0.00, 2421.00, 220.00, 1.00, 1283.00, 52375.00, 12...
## $ state        <chr> "failed", "failed", "failed", "failed", "canceled"...
## $ backers      <dbl> 0, 15, 3, 1, 14, 224, 16, 40, 58, 43, 0, 100, 0, 0...
## $ country      <chr> "GB", "US", "US", "US", "US", "US", "US", "US", "U...
## $ `usd pledged` <dbl> 0.00, 100.00, 220.00, 1.00, 1283.00, 52375.00, 120...
## $ usd_pledged_real <dbl> 0.00, 2421.00, 220.00, 1.00, 1283.00, 52375.00, 12...
## $ usd_goal_real  <dbl> 1533.95, 30000.00, 45000.00, 5000.00, 19500.00, 50...
```

Explorative Data Analysis

Overview of project state:

```
kickstarter %>%
  group_by(state) %>%
  count()
```

```
## # A tibble: 6 x 2
## # Groups:   state [6]
##   state      n
##   <chr>    <int>
## 1 canceled  38779
## 2 failed   197719
## 3 live      2799
## 4 successful 133956
## 5 suspended  1846
## 6 undefined  3562
```

```
kickstarter %>%
  filter(state == "successful") %>%
  summarize(success_rate = n()/nrow(kickstarter))
```

```
## # A tibble: 1 x 1
##   success_rate
##   <dbl>
## 1      0.354
```

Overview of usd dollars pledged and goal:

```
overview_usd_pledged_real <- kickstarter %>%
  summarize(mean_usd_pledged_real = mean(usd_pledged_real),
            median_usd_pledged_real = median(usd_pledged_real),
            sd_usd_pledged_real = sd(usd_pledged_real))

overview_usd_goal_real <- kickstarter %>%
  summarize(mean_usd_goal_real = mean(usd_goal_real),
            median_usd_goal_real = median(usd_goal_real),
            sd_usd_goal_real = sd(usd_goal_real))
```

```
overview_usd_pledged_real
```

```
## # A tibble: 1 x 3
##   mean_usd_pledged_real median_usd_pledged_real sd_usd_pledged_real
##           <dbl>           <dbl>           <dbl>
## 1           9059.           624.           90973.
```

```
overview_usd_goal_real
```

```
## # A tibble: 1 x 3
##   mean_usd_goal_real median_usd_goal_real sd_usd_goal_real
##           <dbl>           <dbl>           <dbl>
## 1          45454.           5500          1152950.
```

Project Goal & Success

Question: Is the project goal (amount of money a project asks for) associated with whether it is successful?

Hypotheses:

- H_0 : There is no relationship between a project's goal and whether a project is successful or not.
- H_1 : There is a relationship between a project's goal and whether a project is successful or not.

Create a new data frame with projects that are no longer live:

```
# Exclude ongoing projects
kickstarter_not_live <- kickstarter %>%
  filter(state != "live")
```

Create a new variable named `usd_goal_real_tier` to classify `usd_goal_real` into tiers:

```
# perhaps need a better way to categorize values in this column
kickstarter_not_live <- kickstarter_not_live %>%
  mutate(usd_goal_real_tier = case_when(
    usd_goal_real < 1000 ~ 1,
    usd_goal_real >= 1000 & usd_goal_real < 5000 ~ 2,
    usd_goal_real >= 5000 & usd_goal_real < 10000 ~ 3,
    usd_goal_real >= 10000 & usd_goal_real < 20000 ~ 4,
    usd_goal_real >= 20000 & usd_goal_real < 100000 ~ 5,
    usd_goal_real >= 100000 & usd_goal_real < 500000 ~ 6,
    usd_goal_real >= 500000 ~ 7,
  ))

kickstarter_not_live %>%
  group_by(usd_goal_real_tier) %>%
  count()
```

```
## # A tibble: 7 x 2
## # Groups:   usd_goal_real_tier [7]
##   usd_goal_real_tier      n
##           <dbl>  <int>
## 1             1  47494
## 2             2 113993
## 3             3  70338
## 4             4  59285
## 5             5  68727
## 6             6  13231
## 7             7   2794
```

Create a new variable for success state

```
kickstarter_not_live <- kickstarter_not_live %>%  
  mutate(success_state = if_else(state == "successful", 1, 0))
```

We'll run a chi-square test at $\alpha = 0.05$.

```
chisq.test(table(kickstarter_not_live$success_state, kickstarter_not_live$usd_goal_real_tier))
```

```
##  
## Pearson's Chi-squared test  
##  
## data:  table(kickstarter_not_live$success_state, kickstarter_not_live$usd_goal_real_tier)  
## X-squared = 19624, df = 6, p-value < 2.2e-16
```

Our test statistic was 19624, which has a chi-square distribution with 6 degree of freedom under H_0 . This corresponds to a p-value less than $2.2e-16$. So our decision is to reject H_0 , and there is sufficient evidence that there is a relationship between a project's goal and a project's state.

Project Goal Amount & Success

Claim: The more money the project asks for, the less successful it will be in terms of getting funding.

Emile and Olivia

- Does the category of project influence it's chance of success?
- The category "Technology" will be the most likely funded type of projects.

limitations: possible crossover in the main_categories that we cannot screen for

questions for Yue: * how to weight each bar graph the same to show proportionality between main categories?

* how not to hard code sum in kickstarter_success_ratio pipeline

```
kickstart_success = kickstarter %>%  
  mutate(  
    project_funding_ratio = pledged/goal,  
    success = case_when(  
      project_funding_ratio >= 1 ~ "achieved",  
      project_funding_ratio < 1 ~ "not achieved"  
    )  
  )
```

```
kickstater_success_ratios = kickstart_success %>%  
  group_by(success, main_category) %>%  
  count(success) %>%  
  mutate(  
    sum = 137042+241619  
  ) %>%  
  mutate(  
    proportion_achieved = n/sum  
  )
```


```
kickstater_success_ratios
```

```
## # A tibble: 30 x 5  
## # Groups:   success, main_category [30]  
##   success main_category      n    sum proportion_achieved  
##   <chr>    <chr>      <int> <dbl>          <dbl>
```

```
## 1 achieved Art      11601 378661      0.0306
## 2 achieved Comics   5875 378661      0.0155
## 3 achieved Crafts   2138 378661      0.00565
## 4 achieved Dance    2340 378661      0.00618
## 5 achieved Design   10870 378661      0.0287
## 6 achieved Fashion   5690 378661      0.0150
## 7 achieved Film & Video 24063 378661      0.0635
## 8 achieved Food     6178 378661      0.0163
## 9 achieved Games    12747 378661      0.0337
## 10 achieved Journalism 1026 378661      0.00271
## # ... with 20 more rows
```

```
ggplot(data = kickstarter_success_ratios, aes(x = success,
                                              y = proportion_achieved,
                                              fill = main_category)) +

  facet_wrap(. ~ main_category) +
  geom_col() +
  theme(legend.position = "none") +
  labs(x = "Success",
       y = "Proportion of successful projects",
       title = "Most Categories Were Not Successful on Average",
       subtitle = "Faceted by Main Category")
```



main_category vs. success

H_0 : There is no relationship between main_category and success.

H_1 : There is a relationship between main_category and success.

Where the $\alpha = 0.05$

```
chisq.test(table(kickstart_success$main_category, kickstart_success$success))
```

```
##
## Pearson's Chi-squared test
##
## data:  table(kickstart_success$main_category, kickstart_success$success)
## X-squared = 17099, df = 14, p-value < 2.2e-16
```

Given that the p-value is less than $\alpha = 0.05$, we can reject the null hypothesis. Thus, there is sufficient evidence to suggest that there is a relationship between main category and success.

Simulation, tech is the most likely to be funded $H_0: p_{s-tech} \leq p_{s-other}$ $H_1: p_{s-tech} > p_{s-other}$

```
set.seed(1)
```

```
kickstart_boot = kickstart_success %>%
  mutate(
    TECH = case_when(
      main_category == "Technology" ~ "TECH",
      main_category != "Technology" ~ "other"
    )
  )
kickstart_boot
```

```
## # A tibble: 378,661 x 18
##       ID name category main_category currency deadline goal
##   <dbl> <chr> <chr>    <chr>      <chr>    <date>    <dbl>
## 1 1.00e9 The ~ Poetry Publishing GBP      2015-10-09 1000
## 2 1.00e9 Gree~ Narrati~ Film & Video USD      2017-11-01 30000
## 3 1.00e9 Wher~ Narrati~ Film & Video USD      2013-02-26 45000
## 4 1.00e9 Tosh~ Music Music USD      2012-04-16 5000
## 5 1.00e9 Comm~ Film & ~ Film & Video USD      2015-08-29 19500
## 6 1.00e9 Mona~ Restaur~ Food USD      2016-04-01 50000
## 7 1.00e9 Supp~ Food Food USD      2014-12-21 1000
## 8 1.00e9 Chas~ Drinks Food USD      2016-03-17 25000
## 9 1.00e9 SPIN~ Product~ Design USD      2014-05-29 125000
## 10 1.00e8 STUD~ Documen~ Film & Video USD      2014-08-10 65000
## # ... with 378,651 more rows, and 11 more variables: launched <dtm>,
## # pledged <dbl>, state <chr>, backers <dbl>, country <chr>, `usd
## # pledged` <dbl>, usd_pledged_real <dbl>, usd_goal_real <dbl>,
## # project_funding_ratio <dbl>, success <chr>, TECH <chr>
```

```
obs_diff = kickstart_boot %>%
  group_by(TECH) %>%
  count(success) %>%
  mutate(prop = n/sum(n)) %>%
  filter(success == "achieved") %>%
  select(TECH, prop)
obs_diff
```

```
## # A tibble: 2 x 2
## # Groups:   TECH [2]
##   TECH prop
##   <chr> <dbl>
## 1 TECH 0.206
```

```
## 2 other 0.377
diff = obs_diff[2,2] - obs_diff[1,2] #take second row and second column - ...
diff
```

```
##      prop
## 1 0.1708602
```

```
tech = kickstart_boot %>%
  filter(TECH == "TECH")
other = kickstart_boot %>%
  filter(TECH != "TECH")

boot_samp <- numeric(500)

for(i in 1:500){
  boot_tech <- tech %>%
    slice(sample(1:nrow(tech), replace = T)) %>%
    count(success) %>%
    mutate(prop = n/sum(n)) %>%
    select(prop) %>%
    slice(1) %>%
    pull()
  boot_other <- other %>%
    slice(sample(1:nrow(other), replace = T)) %>%
    count(success) %>%
    mutate(prop = n/sum(n)) %>%
    select(prop) %>%
    slice(1) %>%
    pull()

  boot_samp[i] <- boot_tech - boot_other
}

#mean of bootstrap samples:
mean_boot = mean(boot_samp)
boot_samp = tibble(boot_samp)

boot_samp <- boot_samp %>%
  mutate(centered = boot_samp + abs(mean_boot))

boot_samp %>%
  mutate(extreme = if_else(centered > abs(diff), 0, 1)
) %>%
  summarize(p_val = mean(extreme))
```


```
## # A tibble: 1 x 1
##   p_val
##   <dbl>
## 1     1
```

```
ggplot(data = boot_samp, aes(x = centered)) +
  geom_histogram(color = "darkblue", fill = "skyblue") +
  labs(x = "centered means",
       y = "count",
       title = "Bootstrap Distribution of Proportions of Success") +
```

```
geom_vline(xintercept = diff$prop, color = "red")
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```

Bootstrap Distribution of Proportions of Success



We have a p-value of 1, which means there is not sufficient evidence to reject the null hypothesis that the proportion of success for main categories other than Technology is the same or greater than the proportion of success for Technology.